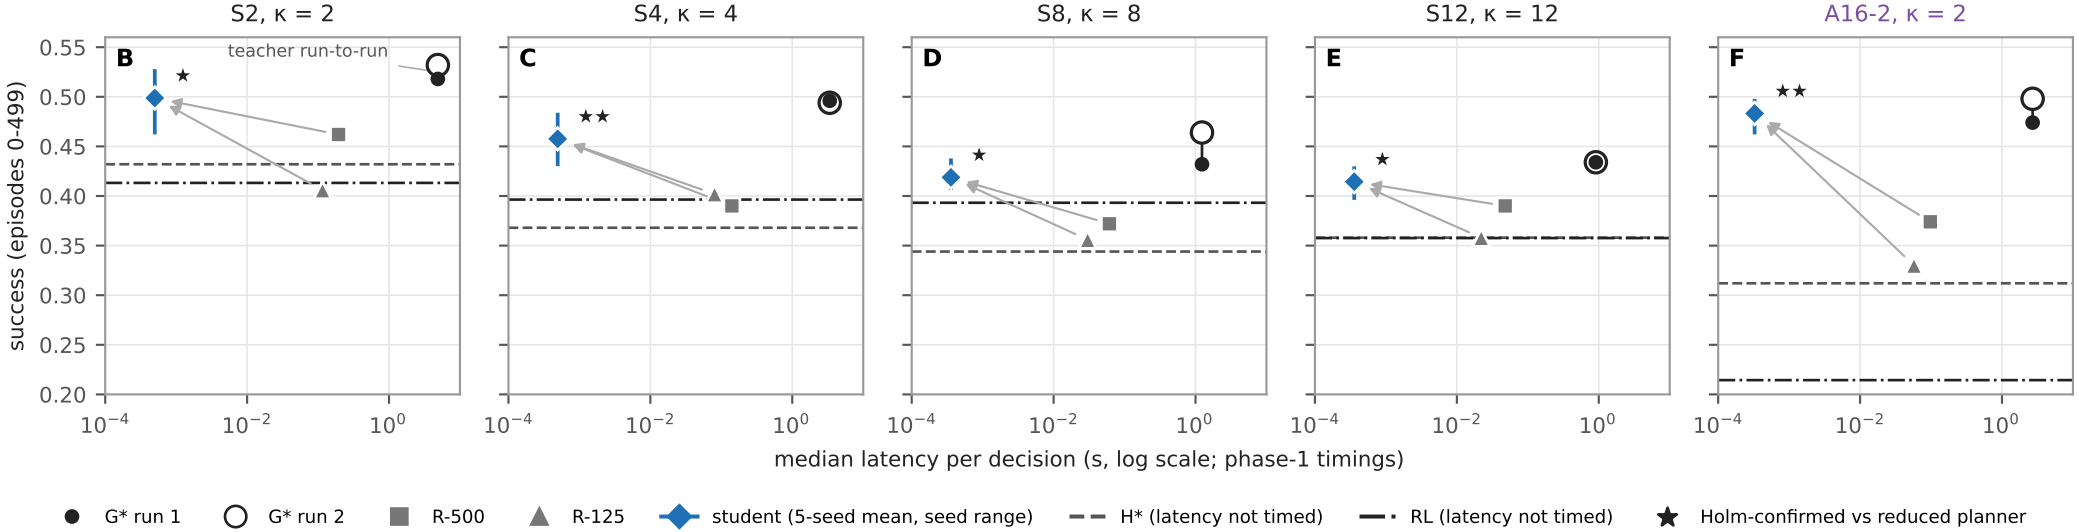
S2, κ = 2
S4, κ = 4
S8, κ = 8
S12, κ = 12
A16-2, κ = 2
B
C
D
E
F
teacher run-to-run
success (episodes 0-499)
0.55
0.50
0.45
0.40
0.35
0.30
0.25
0.20
10^{-4}
10^{-2}
10^{0}
median latency per decision (s, log scale; phase-1 timings)
G* run 1
G* run 2
R-500
R-125
student (5-seed mean, seed range)
H* (latency not timed)
RL (latency not timed)
Holm-confirmed vs reduced planner